

The Self-hosted AI Starter Kit is an open-source template that quickly sets up a local AI environment. Curated by n8n, it provides essential tools for creating secure, self-hosted AI workflows.

[n8n.io](#)

Apache-2.0 license

9.7k stars 2.3k forks Branches Tags Activity

Star

Notifications

Code Pull requests 3 Discussions Actions Security Insights

1 Branch 0 Tags

Go to file

Go to file

Code


 **aporat** fixed AI template gallery link in README.md (#65)
2fe604a · 3 weeks ago

assets	docs: Add gif to readme (#8)	10 months ago
n8n/demo-data	Rename backup to demo-dat...	2 months ago
.env	Use filesystem mode for bin...	2 months ago
LICENSE	Create LICENSE (#5)	10 months ago
README.md	fixed AI template gallery link...	3 weeks ago
docker-compose.yml	Rename backup to demo-dat...	2 months ago

README Apache-2.0 license

Self-hosted AI starter kit

Self-hosted AI Starter Kit is an open-source Docker Compose template designed to swiftly initialize a comprehensive local AI and low-code development environment.



Curated by <https://github.com/n8n-io>, it combines the self-hosted n8n platform with a curated list of compatible AI products and components to quickly get started with building self-hosted AI workflows.

💡 Tip

[Read the announcement](#)

What's included

- ✅ [Self-hosted n8n](#) - Low-code platform with over 400 integrations and advanced AI components
- ✅ [Ollama](#) - Cross-platform LLM platform to install and run the latest local LLMs
- ✅ [Qdrant](#) - Open-source, high performance vector store with an comprehensive API
- ✅ [PostgreSQL](#) - Workhorse of the Data Engineering world, handles large amounts of data safely.

What you can build

- ★ **AI Agents** for scheduling appointments
- ★ **Summarize Company PDFs** securely without data leaks
- ★ **Smarter Slack Bots** for enhanced company communications and IT operations
- ★ **Private Financial Document Analysis** at minimal cost

Installation

Cloning the Repository

```
git clone https://github.com/n8n-io/self-hosted-ai-starter-kit.git
cd self-hosted-ai-starter-kit
```



Running n8n using Docker Compose

For Nvidia GPU users

```
git clone https://github.com/n8n-io/self-hosted-ai-starter-kit.git
cd self-hosted-ai-starter-kit
docker compose --profile gpu-nvidia up
```



Note

If you have not used your Nvidia GPU with Docker before, please follow the [Ollama Docker instructions](#).

For AMD GPU users on Linux

```
git clone https://github.com/n8n-io/self-hosted-ai-starter-kit.git
cd self-hosted-ai-starter-kit
docker compose --profile gpu-amd up
```



For Mac / Apple Silicon users

If you're using a Mac with an M1 or newer processor, you can't expose your GPU to the Docker instance, unfortunately. There are two options in this case:

1. Run the starter kit fully on CPU, like in the section "For everyone else" below
2. Run Ollama on your Mac for faster inference, and connect to that from the n8n instance

If you want to run Ollama on your mac, check the [Ollama homepage](#) for installation instructions, and run the starter kit as follows:

```
git clone https://github.com/n8n-io/self-hosted-ai-starter-kit.git
cd self-hosted-ai-starter-kit
docker compose up
```



For Mac users running OLLAMA locally

If you're running OLLAMA locally on your Mac (not in Docker), you need to modify the OLLAMA_HOST environment variable in the n8n service configuration. Update the x-n8n section in your Docker Compose file as follows:

```
x-n8n: &service-n8n
# ... other configurations ...
environment:
# ... other environment variables ...
- OLLAMA_HOST=host.docker.internal:11434
```



Additionally, after you see "Editor is now accessible via: <http://localhost:5678/>":

1. Head to <http://localhost:5678/home/credentials>
2. Click on "Local Ollama service"
3. Change the base URL to "<http://host.docker.internal:11434/>"

For everyone else

```
git clone https://github.com/n8n-io/self-hosted-ai-starter-kit.git
cd self-hosted-ai-starter-kit
docker compose --profile cpu up
```



Quick start and usage

The core of the Self-hosted AI Starter Kit is a Docker Compose file, pre-configured with network and storage settings, minimizing the need for additional installations. After completing the installation steps above, simply follow the steps below to get started.

1. Open <http://localhost:5678/> in your browser to set up n8n. You'll only have to do this once.
2. Open the included workflow: <http://localhost:5678/workflow/srOnR8PAY3u4RSwb>
3. Click the **Chat** button at the bottom of the canvas, to start running the workflow.
4. If this is the first time you're running the workflow, you may need to wait until Ollama finishes downloading Llama3.2. You can inspect the docker console logs to check on the progress.

To open n8n at any time, visit <http://localhost:5678/> in your browser.

With your n8n instance, you'll have access to over 400 integrations and a suite of basic and advanced AI nodes such as [AI Agent](#), [Text classifier](#), and [Information Extractor](#) nodes. To keep everything local, just remember to use the Ollama node for your language model and Qdrant as your vector store.

Note

This starter kit is designed to help you get started with self-hosted AI workflows. While it's not fully optimized for production environments, it combines robust components that work well together for proof-of-concept projects. You can customize it to meet your specific needs

Upgrading

- **For Nvidia GPU setups:**

```
docker compose --profile gpu-nvidia pull  
docker compose create && docker compose --profile gpu-nvidia up
```



- **For Mac / Apple Silicon users**

```
docker compose pull  
docker compose create && docker compose up
```



- **For Non-GPU setups:**

```
docker compose --profile cpu pull  
docker compose create && docker compose --profile cpu up
```



Recommended reading

n8n is full of useful content for getting started quickly with its AI concepts and nodes. If you run into an issue, go to [support](#).

- [AI agents for developers: from theory to practice with n8n](#)
- [Tutorial: Build an AI workflow in n8n](#)
- [Langchain Concepts in n8n](#)
- [Demonstration of key differences between agents and chains](#)
- [What are vector databases?](#)

Video walkthrough

- [Installing and using Local AI for n8n](#)

More AI templates

For more AI workflow ideas, visit the [official n8n AI template gallery](#). From each workflow, select the **Use workflow** button to automatically import the workflow into your local n8n instance.

Learn AI key concepts

- [AI Agent Chat](#)
- [AI chat with any data source \(using the n8n workflow too\)](#)
- [Chat with OpenAI Assistant \(by adding a memory\)](#)
- [Use an open-source LLM \(via Hugging Face\)](#)
- [Chat with PDF docs using AI \(quoting sources\)](#)
- [AI agent that can scrape webpages](#)

Local AI templates

- [Tax Code Assistant](#)
- [Breakdown Documents into Study Notes with MistralAI and Qdrant](#)
- [Financial Documents Assistant using Qdrant and Mistral.ai](#)
- [Recipe Recommendations with Qdrant and Mistral](#)

Tips & tricks

Accessing local files

The self-hosted AI starter kit will create a shared folder (by default, located in the same directory) which is mounted to the n8n container and allows n8n to access files on disk. This folder within the n8n container is located at `/data/shared` -- this is the path you'll need to use in nodes that interact with the local filesystem.

Nodes that interact with the local filesystem

- [Read/Write Files from Disk](#)
- [Local File Trigger](#)
- [Execute Command](#)



License

Releases

No releases published

Packages

No packages published

Contributors 15

